

Predict Diabetes using Perceptron

Loading the Libraries and Dataset

```
import pandas as pd
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt
```

```
dat_know = pd.read_csv('diabetes.csv')
```

```
dat_know.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
dat_know.shape
```

```
(768, 9)
```

```
dat_know.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Pregnancies	768 non-null	int64

1	Glucose	768	non-null	int64
2	BloodPressure	768	non-null	int64
3	SkinThickness	768	non-null	int64
4	Insulin	768	non-null	int64
5	BMI	768	non-null	float64
6	DiabetesPedigreeFunction	768	non-null	float64
7	Age	768	non-null	int64
8	Outcome	768	non-null	int64


dtypes: float64(2), int64(7)
memory usage: 54.1 KB

Exploratory Data Analysis

dat_know.describe()					
Pregnancies Glucose BloodPressure SkinThickness					
Insulin \					
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000
BMI DiabetesPedigreeFunction Age Outcome					
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958	0.476951
std	7.884160	0.331329	11.760232	0.476951	0.476951
min	0.000000	0.078000	21.000000	0.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000	1.000000

sns.countplot(x = 'Outcome',data = dat_know)

<Axes: xlabel='Outcome', ylabel='count'>



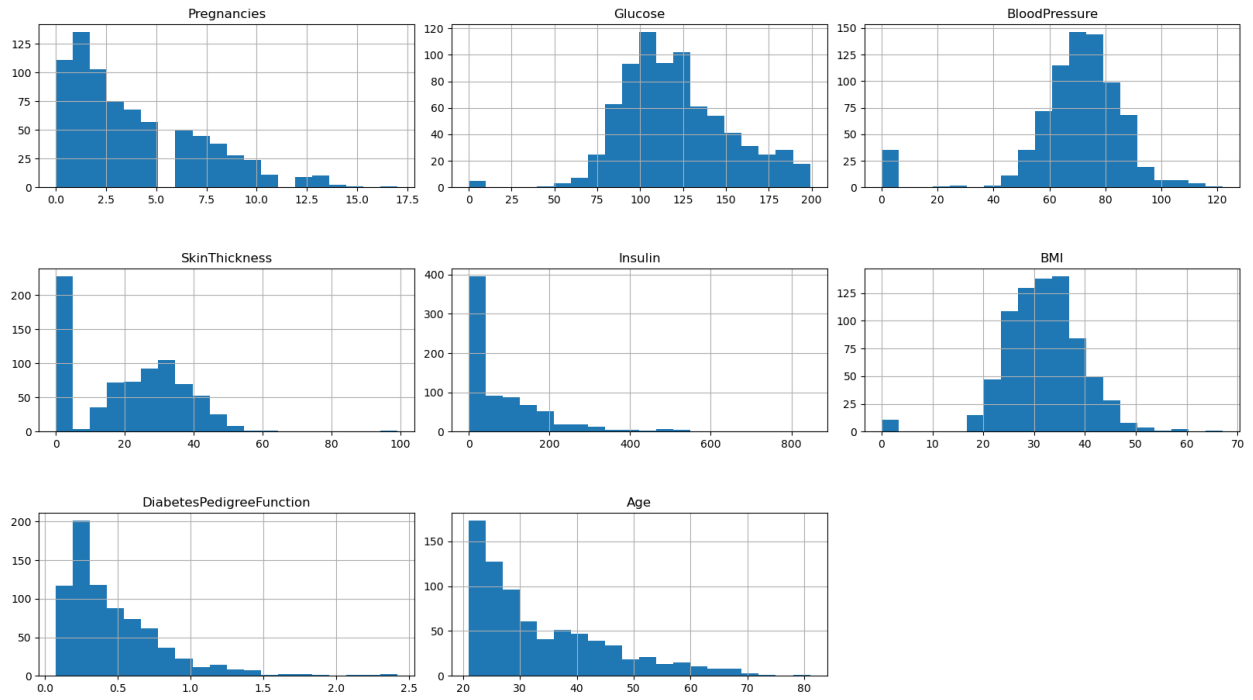
```
import itertools
import math

col = dat_know.columns[:8]
plt.subplots(figsize = (20, 15))
length = len(col)
num_rows = math.ceil(length/2)


for i, j in itertools.zip_longest(col, range(length)):
    plt.subplot(num_rows, 3, j + 1)
    plt.subplots_adjust(wspace = 0.1, hspace = 0.5)
    dat_know[i].hist(bins = 20)
    plt.title(i)
plt.show()
```

C:\Users\Vineeth M R\AppData\Local\Temp\ipykernel_15728\957727032.py:10: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

```
plt.subplot(num_rows, 3, j + 1)
```



```
sns.heatmap(dat_know.corr(), annot = True)  
plt.show()
```



```
dat_new_1 = dat_know

dat_new_1[["Glucose", "BloodPressure", "SkinThickness", "Insulin",
"BMi"]] = dat_new_1[["Glucose", "BloodPressure", "SkinThickness",
"Insulin", "BMi"]].replace(0, np.NaN)

dat_new_1.isnull().sum()

Pregnancies      0
Glucose          5
BloodPressure     35
SkinThickness    227
Insulin          374
BMi              11
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

```

dat_new_1["Glucose"].fillna(dat_new_1["Glucose"].mean(), inplace =
True)
dat_new_1["BloodPressure"].fillna(dat_new_1["BloodPressure"].mean(),
inplace = True)
dat_new_1["SkinThickness"].fillna(dat_new_1["SkinThickness"].mean(),
inplace = True)
dat_new_1["Insulin"].fillna(dat_new_1["Insulin"].mean(), inplace =
True)
dat_new_1["BMI"].fillna(dat_new_1["BMI"].mean(), inplace = True)

dat_new_1.describe()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
768.000000				
mean	3.845052	121.686763	72.405184	29.153420
155.548223				
std	3.369578	30.435949	12.096346	8.790942
85.021108				
min	0.000000	44.000000	24.000000	7.000000
14.000000				
25%	1.000000	99.750000	64.000000	25.000000
121.500000				
50%	3.000000	117.000000	72.202592	29.153420
155.548223				
75%	6.000000	140.250000	80.000000	32.000000
155.548223				
max	17.000000	199.000000	122.000000	99.000000
846.000000				

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.457464	0.471876	33.240885	0.348958
std	6.875151	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.400000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

Training the Dataset

```

from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler(feature_range = (0, 1))
data_scale = sc.fit_transform(dat_new_1)

data_scale = pd.DataFrame(data_scale)

```

```

X = data_scale.iloc[:, [1,2, 4, 5, 7]].values
y = data_scale.iloc[:, 8].values

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.8, random_state = 50, stratify = dat_new_1['Outcome'] )

print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("Y_train shape:", y_train.shape)
print("Y_test shape:", y_test.shape)

X_train shape: (153, 5)
X_test shape: (615, 5)
Y_train shape: (153,)
Y_test shape: (615,)

```

Perceptron Algorithm

```

class Perceptron:
    def __init__(self, learning_rate=0.01, epochs=100):
        self.learning_rate = learning_rate
        self.epochs = epochs

    def fit(self, X, y):
        self.weights = np.zeros(1 + X.shape[1])

        self.errors = []

        for _ in range(self.epochs):
            error = 0
            for xi, target in zip(X, y):
                update = self.learning_rate * (target -
self.predict(xi))
                #print("Ypredict = ", self.predict(xi))
                #print("Y = ", target)
                self.weights[1:] += update * xi
                self.weights[0] += update
                error += int(update != 0.0)

            self.errors.append(error)

        return self
        print(self.errors)
    def net_input(self, X):
        return np.dot(X, self.weights[1:]) + self.weights[0]

    def predict(self, X):
        return np.where(self.net_input(X) >= 0.0, 1, 0)

```

```

def plot_linear_sep(self, X, y):
    plt.figure(figsize=(8, 6))
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired,
marker='o')

    ax = plt.gca()
    xlim = ax.get_xlim()
    ylim = ax.get_ylim()

    xx, yy = np.meshgrid(np.linspace(xlim[0], xlim[1], 50),
                        np.linspace(ylim[0], ylim[1], 50))

    Z = self.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.6)
    plt.xlim(xlim)
    plt.ylim(ylim)

    plt.title('Linearly separable classes')
    plt.xlabel('X')
    plt.ylabel('Y')
    plt.show()

perceptron = Perceptron(learning_rate=0.01, epochs=100)
perceptron.fit(X_train, y_train)

<__main__.Perceptron at 0x199c1171090>

y_pred = perceptron.predict(X_test)

```

Accuracy

```

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")

Accuracy: 76.42%

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm

array([[355, 45],
       [100, 115]], dtype=int64)

from sklearn.datasets import make_classification

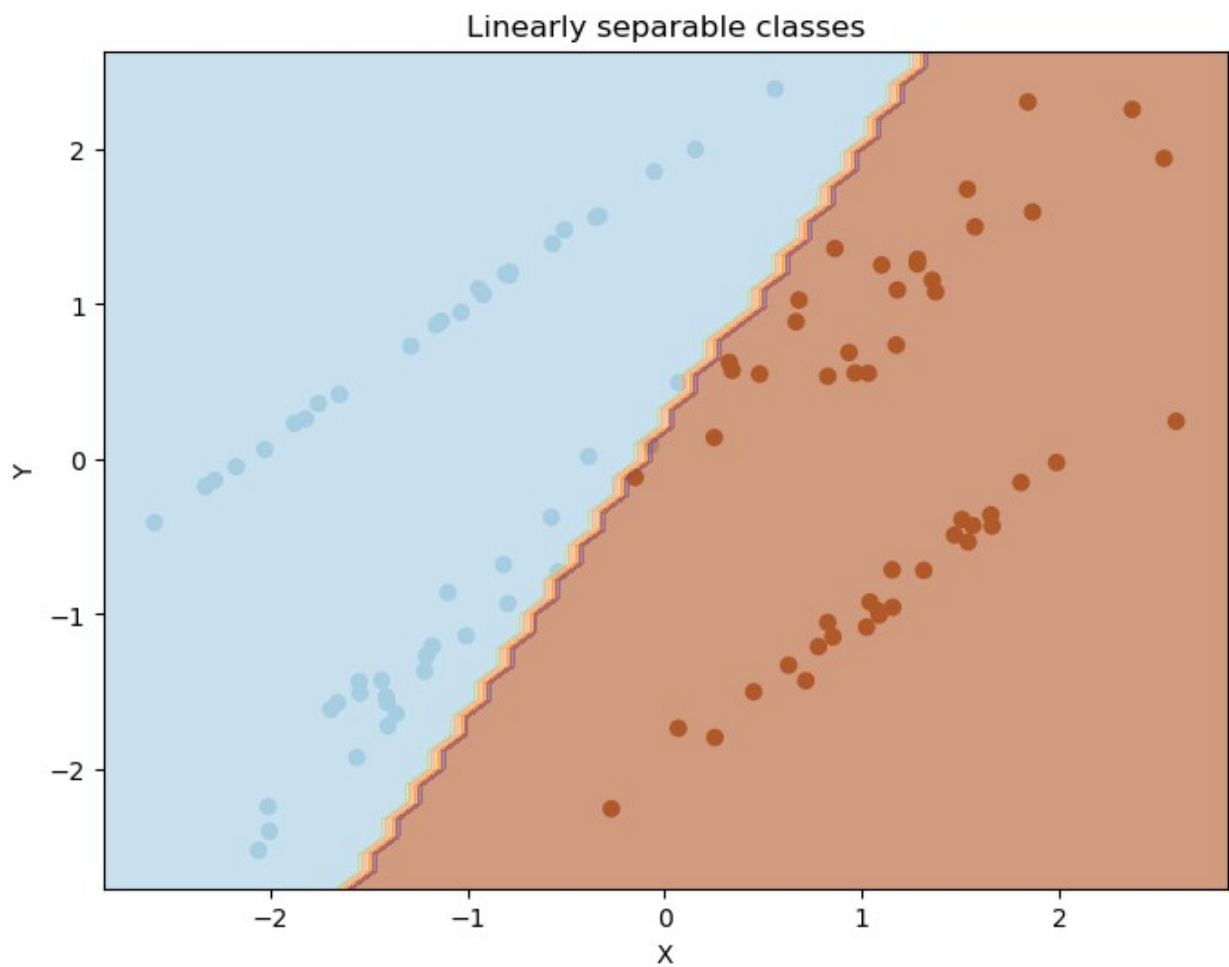
# Generate a synthetic dataset for demonstration
X, y = make_classification(n_samples=100, n_features=2,
n_informative=2, n_redundant=0, random_state=42)

```



```
# Create and train a perceptron model
clf = Perceptron()
clf.fit(X, y)

# Plot the decision boundary
clf.plot_linear_sep(X, y)
```



```
sns.heatmap(pd.DataFrame(cm), annot=True)

<Axes: >
```

